

Tentamen: Statistical Machine Learning (NB054E)

17 January 2017, 0830-1130(1200)

*Write your **name and student number at the top of each sheet**. On each page, indicate page number and total number of pages.*

Please, write clearly! *Make sure to properly motivate all answers, and do not forget to include intermediate steps in your calculations: even if your final answer is wrong, you may still gain some points in that way. You may refer to the Bishop book for relevant equations, etc. One personal “cheat sheet” (a single A4 paper sheet) is allowed.*

Assignment 1

A factory produces thingies. 8% of the thingies is of quality $x = 1$, 32% is of quality $x = 2$, and the remaining 60% is of quality $x = 3$. To assess the quality, two nondestructive tests have been developed that can be performed on a thingy. The first test, Z gives a numerical score z between 0 and 1. The conditional probability density of z depending on the quality x is

$$\begin{aligned}p(z|x=1) &= \alpha_1 z \\p(z|x=2) &= \alpha_2(1-z) \\p(z|x=3) &= \alpha_3(z + \exp(-z))\end{aligned}$$

in which α_i , $i = 1, \dots, 3$ are constants.

Question 1.1 (a) Show that $\alpha_1 = 2$ and $\alpha_2 = 2$. (b) Compute α_3 .

Now there is another test Y . Its outcome is binary $y = 0$ or $y = 1$. The relation of y with the quality is deterministic:

$$\begin{aligned}p(y=1|x=1) &= 1 \\p(y=1|x=2) &= 1 \\p(y=1|x=3) &= 0\end{aligned}$$

Question 1.2 Compute $p(x=i|y=1)$ for $i = 1, \dots, 3$ using Bayes' rule

Tests can be performed repeatedly on a thingy, and test outcomes are independent of each other if performed on the same thingy. Suppose we have performed both test Y and Z once.

Question 1.3 Compute $p(x=i|z, y=1)$ for $i = 1, \dots, 3$

We performed both tests Y and Z once, with result $y = 1$ and $z = 0.5$. We still want to have more certainty about the quality of the thingy. You have the choice to again perform Y , Z , or both.

Question 1.4 Which of these tests, or combination of tests make sense and which not? For the test(s) you choose in this second round, do you expect to find a higher, lower or identical score than the first time? Motivate your answer.

Assignment 2

Consider a stochastic variable k that can have N outcomes $k = 1, \dots, N$. We want to make a probability model $\vec{p} = (p_1, \dots, p_N)$ for this variable, i.e. assign probabilities p_k to each of the possible outcomes k . We do not know \vec{p} , but suppose we do know that the expected value of a certain given function $f(k)$, denoted f_k for short, has the value F , i.e.

$$\langle f_k \rangle_{\vec{p}} \equiv \sum_{k=1}^N p_k f_k = F \quad (1)$$

Unfortunately, just knowing the expectation value $\langle f_k \rangle_{\vec{p}} = F$ is not sufficient to uniquely determine the probabilities \vec{p} .

Question 2.1 Take $N = 3$, $f_k = k$, and $F = 2$, and show by example that there are at least two different distributions $\vec{p}^{(1)}$ and $\vec{p}^{(2)}$ that satisfy (1)

When confronted with a probability distribution in which only a few constraints are known, sometimes the *maximum entropy* (maxent) procedure is used. The entropy is defined as

$$H(\vec{p}) = - \sum_{k=1}^N p_k \ln p_k$$

Question 2.2 Calculate the entropy for the two distributions you provided in question 1.

The idea is that one chooses the distribution that satisfies the constraints, but contains the least information otherwise. In our case, the probability vector is found by maximizing $H(\vec{p})$ under the constraints $\sum_{k=1}^N p_k f_k = F$ and $\sum_{k=1}^N p_k = 1$. For this we write down the Lagrangian with two Lagrange multipliers λ and μ ,

$$L(\vec{p}, \lambda, \mu) = H(\vec{p}) + \lambda \left(\sum_{k=1}^N p_k - 1 \right) + \mu \left(\sum_{k=1}^N p_k f_k - F \right) \quad (2)$$

To maximize L , we first maximize with respect to \vec{p} .

Question 2.3 Set the gradient of the Lagrangian with respect to \vec{p} equal to zero and verify that the probabilities in the maximum satisfy

$$\log(p_k) + 1 + \lambda + \mu f_k = 0 \quad (3)$$

Now λ can be eliminated using the normalisation constraint

Question 2.4 Show, starting from (3) and eliminating λ that the maxent probabilities are of the form

$$p_k = \frac{\exp(-\mu f_k)}{\sum_{j=1}^N \exp(-\mu f_j)} \quad (4)$$

Question 2.5 (a) Argue why, for the given situation (with $N = 3$, $f_k = k$, and $F = 2$), the distribution $\vec{p} = (\frac{1}{3}, \frac{1}{3}, \frac{1}{3})$ is the ‘obvious’ maxent solution; (b) Calculate the maximum entropy solution in case the expectation value is decreased to $F = 1.5$.

THERE ARE TWO MORE QUESTIONS ON PAGE 3 AND 4

Assignment 3

Consider a Gaussian distribution $p(x, y, z)$ with mean

$$\boldsymbol{\mu} = (1, 3, 5)^T \quad (5)$$

and covariance

$$\boldsymbol{\Sigma} = \begin{pmatrix} 2 & 1 & 0 \\ 1 & 2 & 1 \\ 0 & 1 & 4 \end{pmatrix} \quad (6)$$

Since $p(x, y, z)$ is a Gaussian distribution, we know that $p(x, z)$ should also be a Gaussian with a certain mean vector and covariance matrix.

Question 3.1 *What are the mean and covariance of the distribution $p(x, z)$? Sketch a few contours of constant probability.*

Next we look at the probability distribution of $p(x|z)$.

Question 3.2 *(a) Is this a marginal, conditional or joint distribution? Explain. (b) Give an expression for the distribution of $p(x|z)$.*

We now consider a *different* Gaussian distribution $p(u, v, w) = p(u)p(v|u)p(w|u)$, defined in terms of conditional distributions as

$$p(u) = \mathcal{N}(u|\mu_0, \sigma^2) \quad (7)$$

$$p(v|u) = \mathcal{N}(v|c \cdot u, s^2) \quad (8)$$

$$p(w|u) = \mathcal{N}(w|d \cdot u, t^2) \quad (9)$$

with μ_0 , σ^2 , c , s^2 , d and t^2 constant model parameters. The conditional distribution $p(u|v)$ is a Gaussian with a certain mean $\mu_{u|v}$ and variance $\sigma_{u|v}^2$. The question is: what are they?

Question 3.3 *Show that mean $\mu_{u|v}$ and variance $\sigma_{u|v}^2$ of the distribution $p(u|v)$ are given by*

$$\mu_{u|v} = \frac{\frac{\mu_0}{\sigma^2} + \frac{cv}{s^2}}{\frac{1}{\sigma^2} + \frac{c^2}{s^2}} \quad (10)$$

$$\frac{1}{\sigma_{u|v}^2} = \frac{1}{\sigma^2} + \frac{c^2}{s^2} \quad (11)$$

The conditional distribution $p(u|v, w)$ is also a Gaussian.

Question 3.4 *Give an expression for the distribution $p(u|v, w)$ in terms of the model parameters.*

THERE IS ONE MORE QUESTION ON PAGE 4

Assignment 4

In the early 80's Hinton and Sejnowski introduced the Boltzmann machine as a class of recurrent neural networks. Boltzmann machines can be understood as probability distributions of binary variables $\vec{s} = (s_1, \dots, s_N)$, where $s_i = \pm 1$. The Boltzmann machine is parameterized by the matrix W , in which an element w_{ij} represents the connection strength between variables s_i and s_j . The distribution modeled by the Boltzmann machine is defined by

$$P(\vec{s}|W) = \frac{1}{Z(W)} \exp\left(\sum_{i=1}^N \sum_{j=1}^N w_{ij} s_i s_j\right) \quad (12)$$

with $Z(W)$ the normalisation constant,

$$Z(W) = \sum_{s_1=\pm 1} \dots \sum_{s_N=\pm 1} \exp\left(\sum_{i=1}^N \sum_{j=1}^N w_{ij} s_i s_j\right) \quad (13)$$

Question 4.1 Explain how a Boltzmann machine models that certain variables are likely to be both ‘on’ or ‘off’ (+1/−1) together, have opposite values or have no relation at all.

Suppose that we have an i.i.d¹ data set $D = \{\vec{s}^1, \dots, \vec{s}^M\}$. Note that each observation is a vector $\vec{s}^m = (s_1^m, \dots, s_N^m)$. The goal is to learn the parameters in the connectivity matrix W from the data. For that, Hinton & Sejnowski introduced their famous Boltzmann machine learning rule,

$$\Delta w_{ij} = \eta \left[\langle s_i s_j \rangle_{\text{Data}} - \langle s_i s_j \rangle_{P(\vec{s}|W)} \right] \quad (14)$$

So learning is proportional to the difference of two terms. The first term is the empirical correlation between s_i and s_j in the data set:

$$\langle s_i s_j \rangle_{\text{Data}} \equiv \frac{1}{M} \sum_{m=1}^M s_i^m s_j^m \quad (15)$$

The second term is the correlation s_i and s_j under the current model $P(\vec{s}|W)$:

$$\langle s_i s_j \rangle_W \equiv \sum_{s_1=\pm 1} \dots \sum_{s_N=\pm 1} s_i s_j P(\vec{s}|W) \quad (16)$$

We will derive the Boltzmann machine learning rule as gradient ascent on the loglikelihood.

Question 4.2 Show that the log-likelihood of $L(W)$ for the data set D is given by

$$L(W) = M \left(\sum_{i=1}^N \sum_{j=1}^N w_{ij} \langle s_i s_j \rangle_{\text{Data}} - \ln Z(W) \right) \quad (17)$$

with $\langle s_i s_j \rangle_{\text{Data}}$ as defined in (15) and $Z(W)$ as in (13).

Now we have to take the gradient of $\frac{1}{M} L(W)$. Let us first consider the term with $\ln(Z(W))$,

Question 4.3 Show that

$$\frac{\partial}{\partial w_{ij}} \ln Z(W) = \langle s_i s_j \rangle_W$$

with $\langle s_i s_j \rangle_W$ defined as in (16).

Question 4.4 Combine these results to show that the gradient ascent on the log-likelihood (17)

$$\Delta w_{ij} = \frac{\eta}{M} \frac{\partial}{\partial w_{ij}} L(W)$$

leads to the Boltzmann machine learning rule (14).

¹independent and identically distributed